

The Role of AI and Machine Learning in Network Security

Amisha Raje | K. Lahari | Nishita Kamal Lath

April 30, 2025

Word count: 3942

1 Introduction

In our linked world, where cyberthreats are constantly growing in complexity and scope, network security has become more and more important. Modern advanced persistent threats and zero-day attacks are outperforming traditional network security techniques, which rely on static rule sets and signature-based detection (Buczak and Guven, 2016). With the ability to recognize patterns, detect abnormalities, and adjust to emerging threats in ways that traditional systems are unable to, artificial intelligence (AI) and machine learning (ML) have become increasingly potent technologies in recent years to improve network security systems.

This study investigates the relationship between network security and AI/ML technologies, looking at how these computational methods are changing threat detection, prevention, and response systems. A paradigm shift in how businesses safeguard their digital assets has occurred with the incorporation of AI and ML into security frameworks; proactive and predictive security postures have replaced solely reactive ones. (Buczak and Guven, 2016)

Security experts have a difficult time keeping an eye on and safeguarding these technologies as networks get more complicated and threats get more complex. By automating threat identification, decreasing false positives, analyzing massive volumes of network data in real-time, and even anticipating possible vulnerabilities before they can be exploited, AI and ML provide promising answers (Nguyen and Armitage, 2018). These technologies do, however, also bring with them new issues and concerns, such as the necessity for suitable governance structures, ethical ramifications, and possible hostile attacks on ML models themselves.

The purpose of this presentation is to give a thorough understanding of how AI and ML are transforming network security to undergraduate computer science students. We start with an overview of pertinent literature before moving on to basic ideas in AI/ML and network security. Following that,

the paper examines particular applications, real-world case studies, related difficulties, and potential avenues for further study and application.

2 Literature Review

Over the past ten years, there has been a notable increase in research interest in the use of AI and ML techniques in network security. This section examines significant contributions in this quickly developing discipline, emphasizing significant publications and research trends.

By contrasting several classifier methods for network-based attack detection, early research by Tsai et al. (2009) shown the promise of machine learning for intrusion detection systems (IDS). Their study laid the groundwork for future advancements by showing that machine learning techniques might greatly surpass conventional signature-based techniques in detecting fresh assaults.

A thorough analysis of machine learning techniques for cyber security applications, with an emphasis on intrusion detection, was presented by Buczak and Guven (2016). In their work, different machine learning techniques were carefully categorized and their efficacy, computational efficiency, and training needs in security contexts were assessed. Both the potential and constraints of early machine learning applications in security were emphasized in this groundbreaking survey.

With the work of Nguyen and Armitage (2018), who showed that deep neural networks could more successfully mimic the intricate patterns in network traffic linked to sophisticated assaults, deep learning techniques gained popularity. Their method, which made use of neural networks' ability to spot minute patterns in high-dimensional data, demonstrated especially good results in identifying abnormalities that conventional systems would overlook.

Studied a number of feature extraction methods and classification algorithms in order to properly investigate the efficacy of machine learning in malware detection. Their research demonstrated how machine learning techniques may detect malware that had not yet been discovered by extrapolating from known samples, which is a significant benefit above signature-based methods.

The susceptibility of machine learning systems to hostile attacks has been the subject of more recent studies. By modifying inputs in ways that take advantage of the learning algorithms, attackers might purposefully avoid ML-based detection systems, according to Corona, Giacinto and Roli (2013). This brought to light a significant obstacle for implementing machine learning in security environments, as adversaries actively try to get around security mechanisms.

Empirical assessments of machine learning approaches in authentic network environments were carried out by Apruzzese et al. (2018), offering useful insights into their operational efficacy. Their research provided valuable advice on how to choose the best machine learning approaches for par-

ticular security issues and deployment factors in real-world settings.

In order to construct more responsive security systems, Xu et al. (2015) investigated the confluence of machine learning (ML) with security orchestration and automated response. They suggested frameworks for combining threat intelligence with ML-based detection. This line of inquiry has been especially crucial for creating security systems that can identify risks and take appropriate action with little assistance from humans.

As shown by Guo et al. (2018), explainable AI for security applications has become a growing focus of recent research. In order to enable proper human verification and response, their work addressed the crucial necessity for security professionals to comprehend why an ML system classified specific activities as worrisome.

The literature reveals a field in rapid evolution, with research advancing from basic applications of established ML techniques to more sophisticated approaches that address the unique challenges of the security domain. The next sections build upon this foundation to explore how these techniques are applied in practice.

3 Basics of Network Security

The methods, procedures, and policies aimed at safeguarding the confidentiality, accessibility, and integrity of computer networks and data are collectively referred to as network security. Understanding the basic ideas and difficulties in the field of artificial intelligence is crucial before exploring its applications.

3.1 Core Principles of Network Security

Network security is typically built around three primary principles: confidentiality, integrity, and availability—commonly referred to as the CIA triad (?):

- **Confidentiality** ensures that information is accessible only to authorized individuals
- **Integrity** maintains the consistency, accuracy, and trustworthiness of data
- **Availability** guarantees reliable access to information by authorized users

3.2 Traditional Security Mechanisms

Conventional network security approaches include several distinct layers of protection:

Perimeter Security: As the initial line of protection, firewalls, routers, and gateway devices enforce access control policies between networks according to preset rules.

Access Control: Mechanisms for authorization and authentication confirm users' identities and establish what they are allowed to do on the network.

Intrusion Detection and Prevention Systems (IDS/IPS): By keeping an eye out for known attack signatures or suspect trends in network behavior, these systems can either actively block hazardous traffic or generate alarms.

Encryption: Information is protected during storage and transmission via data encryption, which makes sure that even if it is intercepted, it cannot be decrypted without the right keys.

Security Policies and Procedures: Security governance frameworks that specify permissible use, incident response procedures, and security awareness training are examples of non-technical elements.

3.3 Limitations of Traditional Approaches

In spite of their significance, traditional security measures have some drawbacks in the current threat environment:

Reactive Nature: Because signature-based solutions can only detect assaults that have already been reported, networks are exposed to new attack vectors and zero-day exploits.

Rule Maintenance: As networks expand, manually creating and updating security rules becomes more difficult, which frequently results in rule bloat or out-of-date configurations.

False Positives/Negatives: Security staff get alert fatigue as a result of static rule sets' frequent generation of false alerts or failure to detect nuanced threats.

Scalability Challenges: Traditional monitoring approaches struggle to handle the volume, velocity, and variety of data in modern enterprise networks.

Advanced Persistent Threats (APTs): Systems that concentrate on isolated occurrences rather than behavioral patterns may fail to detect sophisticated attacks that evolve over time and use a variety of approaches.

The use of AI and ML techniques, which can learn from data, see trends, and adjust to changing threats in ways that traditional rule-based systems cannot, has flourished as a result of these constraints.

4 Fundamentals of AI and Machine Learning

Gaining a basic understanding of these computational methodologies and their essential security-related techniques is necessary to comprehend how AI and ML improve network security.

4.1 Defining AI and ML

In general, artificial intelligence refers to systems that are capable of carrying out tasks that normally require human intelligence, such as speech recognition, visual perception, language translation, and decision-making. A branch of artificial intelligence called machine learning focuses on algorithms that can learn from data and make judgments or predictions based on it without being specifically programmed for a given task.

4.2 Key ML Approaches in Security Contexts

Supervised Learning: This approach trains models on labeled data, where the correct outputs or classifications are known. In security contexts, supervised learning is commonly used for classifying network traffic as benign or malicious when substantial historical data with known outcomes is available (Buczak and Guven, 2016). Common supervised techniques include:

- Support Vector Machines (SVM): Effective for binary classification problems
- Random Forests: Ensemble methods that build multiple decision trees
- Neural Networks: Multi-layer networks capable of learning complex patterns

Unsupervised Learning: These algorithms identify patterns in data without pre-existing labels. In network security, unsupervised learning excels at anomaly detection, identifying unusual activities that deviate from established normal behavior (Xin et al., 2018). Key techniques include:

- Clustering algorithms (K-means, DBSCAN)
- Dimensionality reduction techniques
- Autoencoders for detecting anomalies

Semi-supervised Learning: A hybrid strategy that combines bigger unlabeled datasets with smaller amounts of annotated data. In security applications where acquiring labeled data is costly or time-consuming, this method is especially beneficial (Mohri, Rostamizadeh and Talwalkar, 2018).

Reinforcement Learning: a paradigm in which agents interact with their surroundings and learn the best practices through trial and error while being rewarded or punished. Reinforcement learning has applications in automated response systems and adaptive defense methods.

4.3 Deep Learning in Security

Neural networks with several hidden layers, which are the hallmark of deep learning, have demonstrated impressive results in security applications:

Convolutional Neural Networks (CNNs): Originally developed for image recognition, CNNs have since been used for security uses such as malware classification using binary representations and network traffic pattern detection.

Recurrent Neural Networks (RNNs) and Long Short-Term Memory Networks (LSTMs): These architectures are particularly effective for sequence data analysis, making them valuable for analyzing network packet sequences and identifying temporal attack patterns (?).

Deep Belief Networks: The ability of these generative models to learn to probabilistically recreate their inputs makes them useful for spotting anomalies in network activity.

4.4 Feature Engineering and Selection

The effectiveness of machine learning in security depends on identifying and extracting relevant features from raw security data.

Network Features: Include flow statistics, protocol types, packet sizes, and inter-arrival delays.

Host-based Features: System calls, process information, file access patterns, and registry changes.

User Behavior Features: Command sequences, resource access times, and login patterns.

In order to improve model performance and computing efficiency—two crucial factors in real-time security monitoring systems—feature selection approaches assist in identifying the most useful qualities while minimizing dimensionality (Mohri, Rostamizadeh and Talwalkar, 2018).

The next section explores how these AI and ML approaches are specifically applied to address network security challenges.

5 Applications of AI/ML in Network Security

Many useful applications throughout the security lifecycle have resulted from the use of AI and ML approaches into network security. The main areas where these technologies have proven especially useful are examined in this section.

5.1 Intrusion Detection Systems (IDS)

One of the most developed uses of machine learning in network security is intrusion detection. By learning the traits of typical network behavior and highlighting deviations, machine learning

(ML)-powered intrusion detection systems (IDS) are able to detect new threats, in contrast to classic signature-based systems (Buczak and Guven, 2016).

Anomaly Detection: Algorithms for semi-supervised and unsupervised learning create baseline models of typical network behavior. The system signals possible intrusions when traffic diverges substantially from these baselines. This method works especially well against zero-day assaults without known signatures (Xin et al., 2018).

Misuse Detection: Supervised learning techniques train on labeled datasets of known attack patterns, allowing systems to identify variations of known attacks even when they don't match exact signatures (Tsai et al., 2009).

Hybrid Approaches: In order to reduce false positives and preserve sensitivity to emerging risks, many contemporary systems integrate anomaly and abuse detection, utilizing the advantages of both techniques (Buczak and Guven, 2016).

5.2 Malware Detection and Classification

ML has transformed malware analysis from signature-based approaches to behavior-based identification (?):

Static Analysis: ML models analyze executable files without running them, examining features like byte sequences, opcodes, API calls, and structural elements to identify malicious characteristics.

Dynamic Analysis: Behavioral features are extracted as malware executes in controlled environments, with ML models identifying suspicious activities like unusual system calls, network connections, or file operations.

Multi-stage Detection: Fast initial screening and more in-depth secondary analysis are frequently combined in modern methods to maximize the trade-off between detection speed and accuracy (?).

5.3 Network Traffic Analysis

Beyond explicit attack detection, ML enhances broader network traffic monitoring:

Traffic Classification: More detailed security controls and enhanced visibility are made possible by ML algorithms ability to classify network traffic by application, protocol, or user (Nguyen and Armitage, 2018).

Encrypted Traffic Analysis: Even when packet contents are encrypted, ML can analyze metadata (packet sizes, timing, destination) to identify potential threats without decrypting traffic—critical for maintaining privacy while ensuring security (Anderson and McGrew, 2018).

Data Exfiltration Detection: ML models detect subtle patterns indicative of data theft, identifying unusual data transfers even when attackers attempt to disguise them as normal traffic.

5.4 Automated Response and Orchestration

Beyond detection, ML increasingly enables automated response to security incidents:

Threat Prioritization: ML algorithms assess multiple factors to determine the severity and urgency of threats, helping security teams focus on the most critical issues first.

Response Recommendation: Systems suggest appropriate countermeasures based on threat characteristics and historical response effectiveness.

Autonomous Mitigation: System response times can sometimes be shortened from hours to seconds by automatically implementing protective measures like blocking malicious IP addresses or isolating compromised systems (Xu et al., 2015).

5.5 Threat Intelligence and Prediction

ML enhances proactive security through predictive capabilities:

Threat Hunting: ML guides security analysts toward potential compromise indicators that warrant further investigation.

Vulnerability Prediction: Models analyze code and system configurations to identify potential security weaknesses before they can be exploited.

Attack Prediction: Analysis of threat intelligence and attacker behavior patterns helps predict likely targets and methods, enabling preemptive defensive measures (Buczak and Guven, 2016).

These applications demonstrate how AI and ML are transforming network security from reactive to proactive, from manual to automated, and from rule-based to intelligence-driven approaches.

6 Real-World Case Studies

Analyzing actual implementations offers important insights into how businesses have effectively used AI and ML for network security. Selected case studies illustrating various applications and their results are presented in this section.

6.1 Large Financial Institution: ML-Based Fraud Detection

An ML-based system was put in place by a global banking company to detect fraudulent transactions and patterns of network access (Apruzzese et al., 2018). The system employed:

Implementation: A three-step process that combines supervised classification for in-depth examination of transactions that have been flagged with unsupervised anomaly detection (using autoencoders) for preliminary screening.

Results: The system reduced false positives by 60% compared to rule-based approaches while simultaneously increasing fraud detection rates by 35%. It also identified several previously undetected attack patterns involving coordinated activities across multiple accounts.

Challenges: Much fine-tuning was necessary for the initial deployment in order to lower false positives without sacrificing detection sensitivity. Clear procedures for analyst review of ML-flagged activities also have to be developed by the organization.

6.2 Telecommunications Provider: Network Traffic Analysis

A major telecommunications provider deployed deep learning models to analyze encrypted network traffic for security threats without decryption (Anderson and McGrew, 2018):

Implementation: The system identified fraudulent traffic patterns while protecting client privacy by analyzing packet metadata, timing, and flow characteristics using convolutional neural networks (CNNs).

Results: Compared to conventional methods that either needed decryption (compromising privacy) or had lower detection rates, the method detected 91% of malware communications and command-and-control traffic without requiring packet decoding.

Challenges: As both typical traffic patterns and attack strategies changed, it became necessary to continuously retrain the model in order to maintain accuracy. The business put in place an automated retraining pipeline that included edge case evaluation by humans.

6.3 Critical Infrastructure: IDS for Industrial Control Systems

An ML-based intrusion detection system created especially for industrial control systems (ICS) was put into place by a power distribution firm (Guo et al., 2018):

Implementation: An LSTM-based anomaly detection model trained on typical operational patterns was integrated with conventional rule-based detection in the system. The addition of explainable AI components, which gave alerts human-interpretable justifications, was a significant innovation.

Results: The system effectively identified a number of complex attacks that imitated valid commands but would have disrupted operations over the course of a 12-month operational period. By assisting analysts in rapidly comprehending warning settings, the explainability features decreased incident reaction time by 47%.

Challenges: The highly specialized nature of ICS environments required extensive domain knowledge integration. In order to capture the distinctive features of industrial procedures and operating states, engineers created domain-specific feature extraction.

6.4 E-commerce Platform: Bot Detection and Prevention

To differentiate between real consumers and synthetic bots engaging in inventory scraping or credential stuffing attacks, a major e-commerce company used machine learning models (Xin et al., 2018):

Implementation: The system used contextual elements, request timing analysis, and behavioral biometrics (keystroke patterns, mouse movements) in an ensemble method. Components of reinforcement learning were constantly modified to accommodate changing bot behaviors.

Results: The accuracy of bot detection increased from 82% using conventional techniques to 97% using the machine learning technology. By stopping bot-initiated sessions from moving through resource-intensive application components, the solution decreased the waste of computing resources.

Challenges: Advanced robots started trying to imitate human behavior. In response, the business added purposeful random components to the system’s challenge-response processes and used multimodal detection techniques.

Organization Type	ML Technique	Security Application	Improvement
Financial Institution	Autoencoders + Classification	Fraud Detection	60% fewer false positives
Telecommunications	CNN	Encrypted Traffic Analysis	91% detection without decryption
Power Distribution	LSTM + Explainable AI	ICS Intrusion Detection	47% faster incident response
E-commerce	Ensemble + RL	Bot Detection	97% detection accuracy

Table 1: Summary of ML implementations across different sectors

These case studies demonstrate the notable advantages and implementation difficulties of machine learning (ML)-based security solutions across various industries and application scenarios.

7 Challenges and Ethical Considerations

Although AI and ML provide strong network security capabilities, their application poses substantial difficulties and brings up crucial ethical issues that need to be resolved for responsible deployment.

7.1 Technical Challenges

Adversarial Machine Learning: The possibility of attacks on ML systems themselves is arguably the most worrisome obstacle. Adversarial examples—inputs specifically crafted to fool ML models—can cause security systems to misclassify malicious activities as benign (Corona, Giacinto and Roli, 2013).

Although they are still being researched, strategies like defensive distillation and adversarial training can increase robustness.

Data Quality and Availability: ML models are only as good as their training data. Security-relevant training datasets often suffer from:

- Class imbalance (rare attack instances compared to normal traffic)
- Mislabeled samples due to undetected past breaches
- Rapidly evolving attack techniques that create training-deployment gaps

Explainability and Interpretability: Many powerful machine learning techniques, especially deep learning methods, operate as "black boxes" with little information about how they make decisions. Security analysts face difficulties because of this opacity since they must comprehend the reasons behind alert generation in order to decide on the best course of action (Guo et al., 2018).

7.2 Operational Challenges

Skills Gap: System design, implementation, and maintenance are made more difficult by organizations' lack of staff with knowledge of both security and machine learning.

Integration with Existing Systems: ML-based solutions usually need to work with and integrate with conventional security infrastructure, which calls for meticulous process modifications and architectural planning.

Continuous Learning Requirements: Attack methods and typical network behavior are always changing, so security is never static. For ML systems to remain effective, frequent retraining and validation are necessary, which adds significant operational overhead.

Alert Management: More alarms can be produced by even ML-enhanced systems than analysts can handle. Effective triage procedures must be put in place by organizations in order to concentrate human resources where they are most required.

7.3 Ethical Considerations

Privacy Implications: ML systems sometimes call for a large amount of data to be collected and analyzed, possibly including sensitive communications or personally identifying information. Businesses must weigh the advantages of security against the effects on privacy (?).

Bias and Fairness: Biases in training data may be unintentionally incorporated into ML systems, which could result in excessive flagging of the behaviors of particular user groups or unfair security enforcement.

Transparency and Consent: Users whose actions are being tracked have a right to know how their information is used and what choices might be made using machine learning.

Accountability: The increasing automation of security choices raises concerns regarding accountability in the event of system failures or false positives that affect authorized users.

8 Future Directions

The use of AI and ML in network security is still developing quickly. This section examines new developments and lines of inquiry that are anticipated to influence the discipline in the years to come.

8.1 Advanced Defensive Techniques

Self-Healing Networks: In addition to detecting attacks, research is moving toward fully autonomous security systems that can isolate vulnerable segments, apply patches, and change network designs automatically without the need for human participation (Xu et al., 2015).

Federated Learning for Security: This method could transform threat intelligence sharing while maintaining privacy and confidentiality by enabling firms to work together to train models without exchanging sensitive data (Mohri, Rostamizadeh and Talwalkar, 2018).

Robust AI against Adversarial Attacks: Research into robust machine learning models that are resistant to manipulation is accelerating as adversarial machine learning advances. Methods such as defensive distillation and guaranteed robustness show promise in building models that are resistant to evasion efforts (Corona, Giacinto and Roli, 2013).

Zero-trust Architecture Integration: AI/ML systems, which continually validate identities and authorizations based on behavioral patterns rather than one-time authentication, will progressively take center stage in zero-trust security frameworks.

8.2 Enhanced Understanding and Context

Explainable AI (XAI) for Security: The development of security technologies that can not only identify risks but also clearly explain their logic to analysts will be facilitated by ongoing research into making machine learning systems more transparent and interpretable (Guo et al., 2018).

Multi-modal Analysis: In order to provide a more comprehensive contextual knowledge of security events, future systems will increasingly incorporate data from a variety of sources, including network traffic, endpoint telemetry, threat information, and user behavior.

Natural Language Processing for Threat Intelligence: Automated threat information extraction and analysis from unstructured sources, such as security blogs, research papers, and dark web forums, will be enhanced by NLP approaches.

8.3 Emerging Application Domains

IoT Security: ML techniques created especially for resource-constrained situations and distinct IoT protocols will become more and more significant as the number of Internet of Things devices increases (Xin et al., 2018).

5G/6G Network Security: The complexity, virtualization, and dynamic nature of next-generation telecommunications networks will necessitate AI-driven security.

Supply Chain Security: The use of machine learning algorithms to identify counterfeiting, tampering, or compromise in supply chains for software and hardware will grow.

Quantum-Safe Security: ML may be useful in identifying threats and modifying security protocols in post-quantum security frameworks as breakthroughs in quantum computing pose a threat to conventional encryption techniques.

8.4 Human-AI Collaboration

Augmented Intelligence: Advanced systems will work to enhance human capabilities rather than replace security analysts, managing repetitive chores and revealing insights that call for human judgment.

Adaptive User Interfaces: Security consoles will progressively adjust to the workflows and skill levels of individual analysts, displaying machine learning findings in ways that are most suited for human comprehension and judgment.

Security Automation: For complicated scenarios, orchestration frameworks will develop to smoothly combine human decision points with automated ML-driven reactions.

9 Conclusion

The integration of artificial intelligence and machine learning into network security represents a significant evolution in how organizations defend their digital assets. These technologies have strong capabilities to identify, stop, and address security risks in ways that are unmatched by conventional methods, as this paper has examined.

The core advantages of AI/ML approaches in security contexts include:

- Ability to identify novel threats without pre-existing signatures
- Capacity to process and find patterns in vast quantities of security data
- Adaptive capabilities that evolve as both normal network behavior and attack techniques change
- Potential for automation that reduces response times and analyst workload

However, the implementation of these technologies is not without challenges. Technical hurdles like adversarial attacks against ML models themselves, data quality issues, and explainability concerns must be addressed. The ethical issues of privacy, bias, accountability, and transparency that come up during the implementation of automated security systems are equally significant.

The case studies presented demonstrate that organizations across sectors are finding ways to successfully navigate these challenges and realize significant security benefits. Their experiences demonstrate the value of careful implementation strategies that take into account operational procedures, ethical considerations, and technical aspects.

In the future, the field will continue to develop quickly. More efficient security features are anticipated as a result of developments in fields like explainable AI, federated learning, and human-AI collaboration. Meanwhile, new application areas like supply chain security and IoT security will open up new possibilities for AI/ML techniques.

The nexus of network security and AI/ML presents abundant research and real-world application opportunities for incoming computer science students. Technical expertise in these fields is necessary for success, but so is an awareness of the larger framework in which security systems function, including organizational, moral, and legal factors.

Intelligent systems that can recognize and react to threats on their own while collaborating with human professionals will probably determine the future of network security. The security community can remain ahead of changing threats while upholding the privacy and trust that are crucial in our globalized society by carefully creating and using these solutions.

References

- Anderson, Blake and David McGrew. 2018. Machine learning for encrypted malware traffic classification: Accounting for noisy labels and non-stationarity. In *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. ACM pp. 1723–1732.
- Apruzzese, Giovanni, Michele Colajanni, Luca Ferretti, Alessandro Guido and Mirco Marchetti. 2018. On the effectiveness of machine and deep learning for cyber security. In *2018 10th International Conference on Cyber Conflict (CyCon)*. IEEE pp. 371–390.
- Buczak, Anna L and Erhan Guven. 2016. “A survey of data mining and machine learning methods for cyber security intrusion detection.” *IEEE Communications Surveys & Tutorials* 18(2):1153–1176.
- Corona, Igino, Giorgio Giacinto and Fabio Roli. 2013. “Adversarial attacks against intrusion detection systems: Taxonomy, solutions and open issues.” *Information Sciences* 239:201–225.
- Guo, Wenbo, Dongliang Mu, Jun Xu, Purui Su, Gang Wang and Xinyu Xing. 2018. LEMNA: Explaining deep learning based security applications. In *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security*. ACM pp. 364–379.
- Mohri, Mehryar, Afshin Rostamizadeh and Ameet Talwalkar. 2018. *Foundations of machine learning*. MIT press.
- Nguyen, Thuy TT and Grenville Armitage. 2018. “A survey of techniques for internet traffic classification using machine learning.” *IEEE Communications Surveys & Tutorials* 10(4):56–76.
- Tsai, Chih-Fong, Yu-Feng Hsu, Chia-Ying Lin and Wei-Yang Lin. 2009. Intrusion detection by machine learning: A review. Vol. 36 Elsevier pp. 11994–12000.
- Xin, Yang, Lingshuang Kong, Zhi Liu, Yuling Chen, Yanmiao Li, Hongliang Zhu, Mingcheng Gao, Haixia Hou and Chunhua Wang. 2018. “Machine learning and deep learning methods for cyber-security.” *IEEE Access* 6:35365–35381.
- Xu, Zhaoyan, Sanjay Ray, Pramod Subramanyan and Sharad Malik. 2015. A unified framework for ML-based unified malware detection. In *2017 IEEE Cybersecurity Development (SecDev)*. IEEE pp. 98–105.